


Real-time vehicle and pedestrian detection

Developing a **robust, real-time detection system** to improve road safety and reduce pedestrian-vehicle collisions in dense urban environments. System targets **live alerts, accurate localization**, and lightweight deployment for in-vehicle and roadside units.



Problem background

The system has four core goals:

- **Enhance driver awareness** by detecting pedestrians and nearby vehicles in real time.
- **Reduce accident risk** via **distance-based warnings** and prioritized alerts.
- Deliver a lightweight, **deployable solution** using deep learning optimized for real-time inference.
- Provide a **foundation for ITS research and autonomous driving assistance modules**.

Baseline approach

Design and implement a real-time detection pipeline combining:



Detector

YOLOv9 for high-speed object detection with class and bbox outputs.



Distance Estimation

Monocular pinhole model for per-object distance approximations.



Alerts

Distance thresholds generate prioritized driver warnings and logging.



Baseline model: YOLOv9 — do we really need to fine-tune it?



Domain adaptation

Learn road geometry, reflections, and low-light artifacts present in dashcam footage.



Class specialization

Focus on relevant road classes (person, car, bus, truck, bike, motor) to reduce confusion.

Fine-tuning benefits — continued

Localization gains

Improved bounding-box regression for small, distant, or partially occluded pedestrians — better distance inputs.

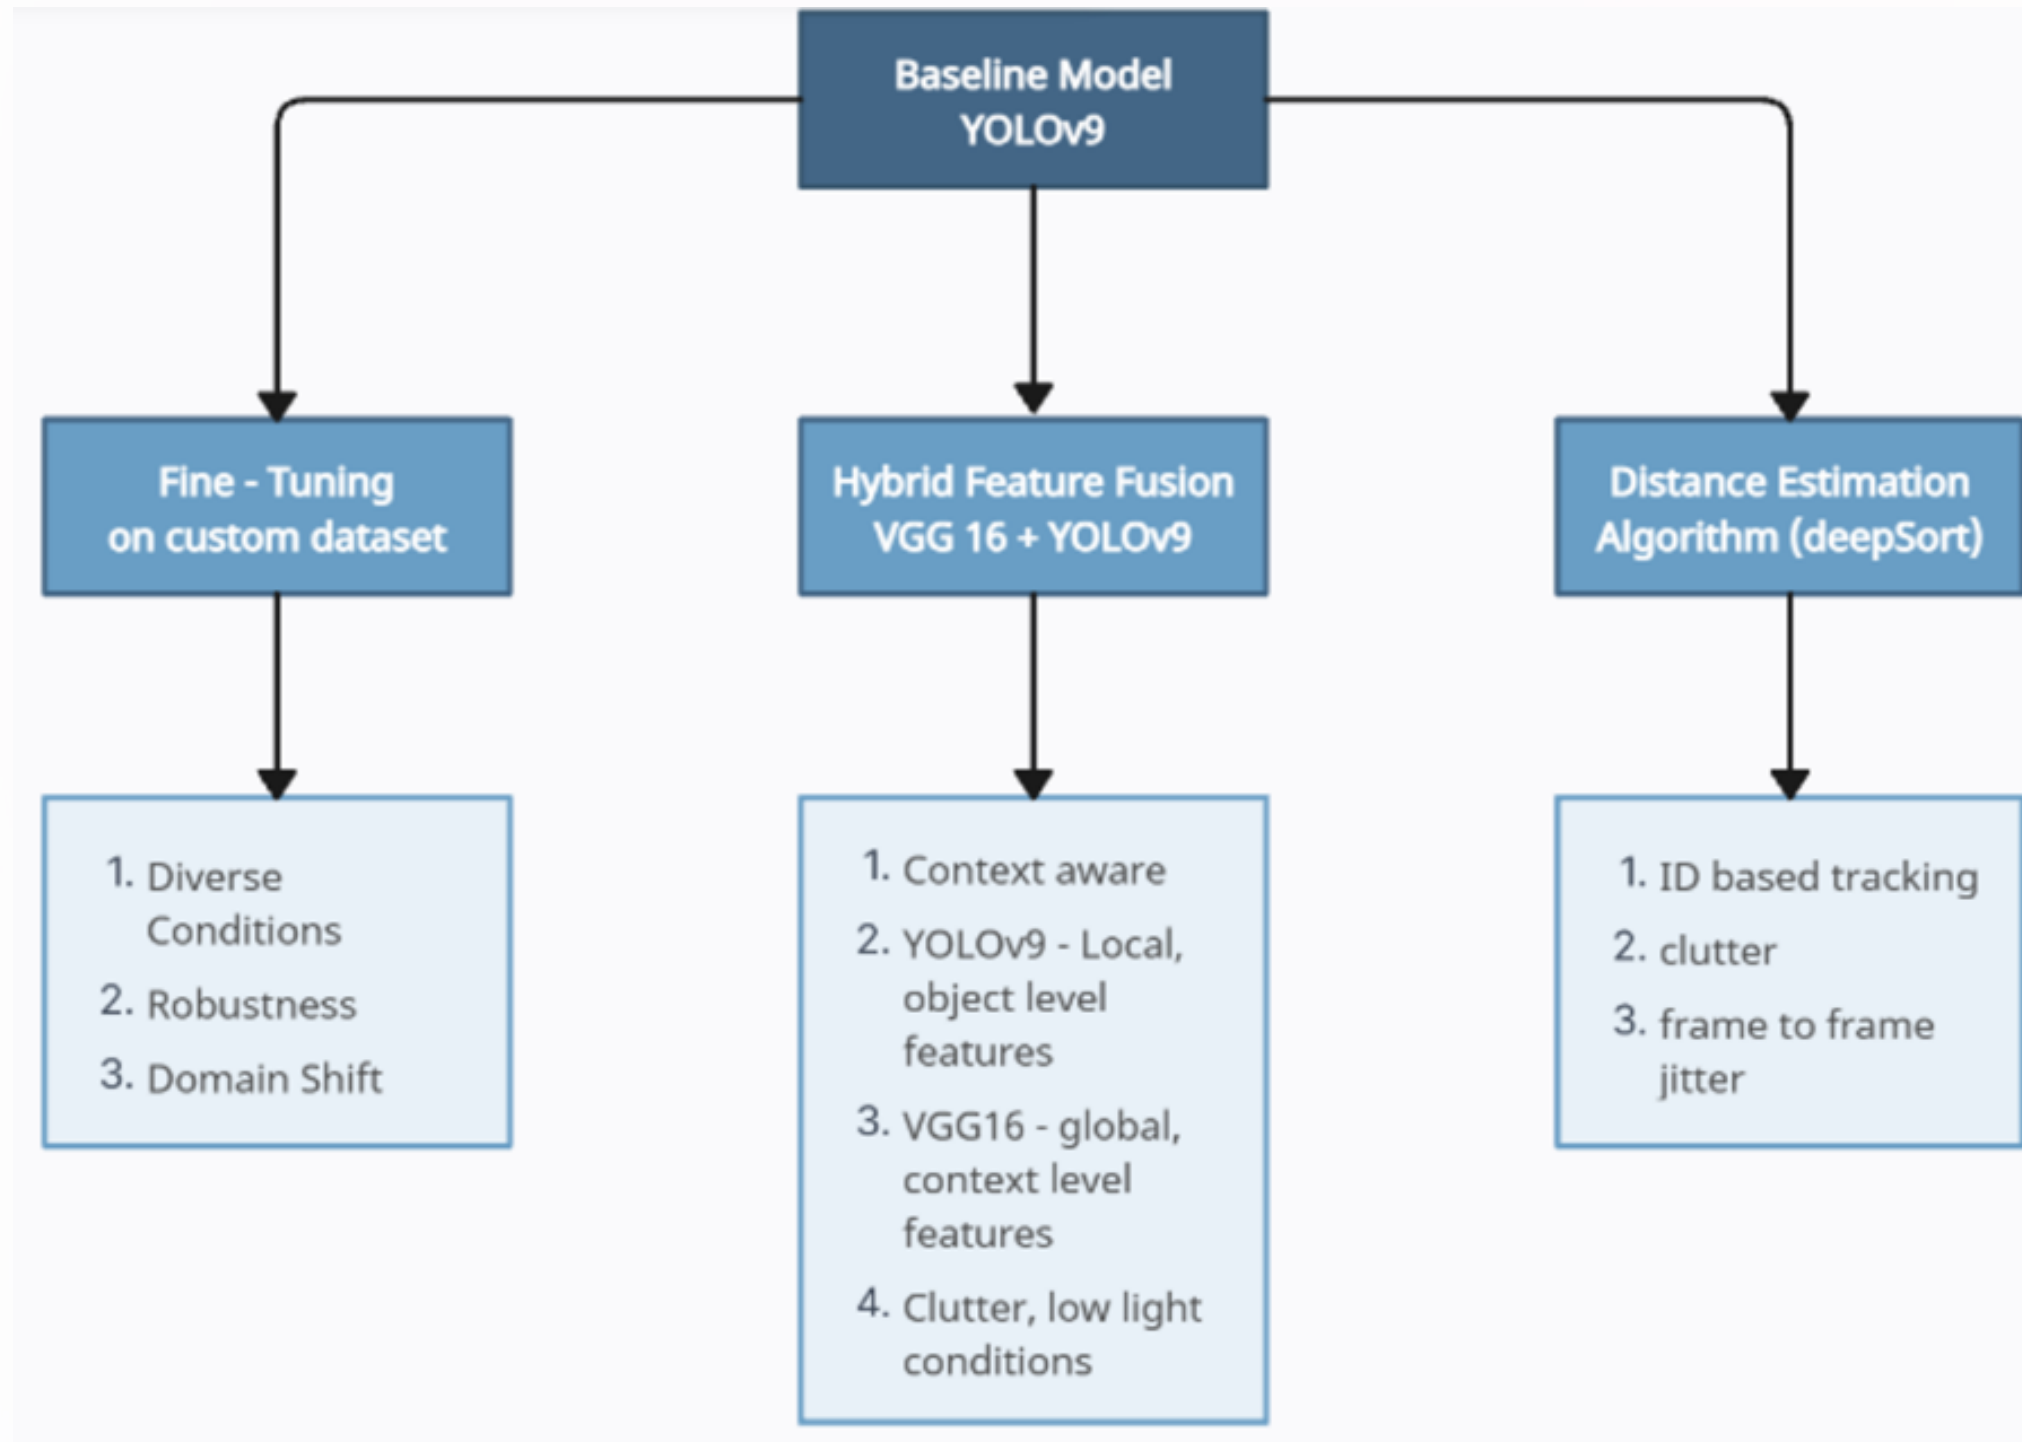
Reduced domain shift

Models tuned on BDD data generalize more consistently across weather, time-of-day, and city scenes.

Faster adaptation

Road-aware weights speed up transfer learning for lane detection, signage, and tracking tasks.

3 – Phase Approach for improving the Baseline Yolov9 model



Fine tuning Yolov9 using the BDD10k Dataset

FineTuning_YOLOv9.ipynb

File Edit View Insert Runtime Tools Help

mands + Code + Text ▶ Run all

Reconnect T4

Share

G

```
# Train
results = model.train(
    data='/content/bdd10k/data.yaml',
    epochs=30,
    imgsz=416,
    batch=32,
    device=0,
    project='runs/train',
    name='bdd10k_fast',
    cache=True,
    patience=10
)
```

	all	Images	Instances	Box(P)	R	mAP50	mAP50-95)
Epoch 23/30	all	2008	36289	0.65	0.469	0.509	0.259
	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size	
	8.96G	1.21	0.801	0.9818	231	416: 100%	247/247 1.4it/s 2:58
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100%
	all	2008	36289	0.641	0.481	0.515	0.263
Epoch 24/30	all	2008	36289	0.641	0.481	0.515	0.263
	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size	
	8.87G	1.203	0.7925	0.9805	277	416: 100%	247/247 1.4it/s 2:58
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100%
	all	2008	36289	0.641	0.481	0.515	0.263

ariables Terminal

21:55

Fine tuned Model Performance:

YOLOv9: mAP50 = **53.7%** on BDD10K dataset

Baseline YOLOv9: mAP50 = **~67%** on COCO

Important Context:

Lower mAP50 compared to COCO baseline is **expected and normal** because:

- 1. **BDD10K is harder** - real-world driving scenarios with occlusions, weather conditions, varying distances
- 2. **Different dataset** - COCO has cleaner, more varied training data
- 3. **Specific domain** - Traffic detection is more challenging than general object detection
- 4. **BDD10k's 5 classes** vs COCO's 80 classes

Validating /content/yolov9/runs/train/bdd10k_fast/weights/best.pt...

Ultralytics 8.3.214 Python-3.12.12 torch-2.8.0+cu126 CUDA:0 (Tesla T4, 15095MiB)

YOLOv9c summary (fused): 156 layers, 25,322,332 parameters, 0 gradients, 102.3 GFLOPs

Class	Images	Instances	Box(P	R	mAP50	mAP50-95)
all	2008	36289	0.672	0.491	0.537	0.276
car	1987	20806	0.739	0.672	0.725	0.46
pedestrian	663	2880	0.592	0.432	0.462	0.214
traffic light	1130	5480	0.67	0.373	0.422	0.152
traffic sign	1643	7123	0.686	0.486	0.539	0.278

Speed: 0.1ms preprocess, 6.7ms inference, 0.0ms loss, 1.2ms postprocess per image

Results saved to /content/yolov9/runs/train/bdd10k_fast


Dataset: BDD100K & BDD10K

Berkely DeepDrive datasets — curated for realistic driving scenarios and multi-task perception.



BDD100K

100K dashcam videos across cities, times, and weather; supports detection, segmentation, and tracking.



BDD10K

10K-image subset for practical fine-tuning and validation with consistent annotation schema.

Distance estimation algorithms

Current method: monocular pinhole camera model. Assumptions and tradeoffs:

- **KNOWN WIDTH:** real-world width of the object (meters).
- **FOCAL LENGTH:** camera focal length in pixels.
- **BOX WIDTH:** width of the detected object's bounding box in pixels.

$$\text{Distance} = \frac{\text{Known Width} \times \text{Focal Length}}{\text{Apparent Width in Pixels}}$$

1. **Assumptions:** known real-world width, object roughly perpendicular to camera, lens distortion is already corrected.
2. **Pros:** simple, fast, easy to implement on edge hardware.
3. **Cons:** sensitive to pose, occlusion, and inter-subject width variation (pedestrian postures).

Algorithm — deepSORT tracker

Why deepSORT is a pragmatic choice for ITS:



Consistent tracking

Maintains identity across frames enabling stable distance and speed estimates.



Real-time ready

Efficient Kalman filter + lightweight CNN embeddings suitable for live monitoring.



Occlusion handling

Re-identifies objects after short losses using motion + appearance features.



Enables analytics

Supports trajectory, speed profiling, safety-zone violation detection, and per-object smoothing.

Next steps — Improvements

1. Integrate **front-end UI and back-end pipeline** to stabilize frame ingestion and buffer handling, **improving the speed and efficiency**
2. Benchmark **alternate distance estimators** like deepSORT and quantify error vs range.
3. Iterate **3 - phase plan for improving the model accuracy** over and above the baseline model

